

Ryan W. Lazar

Data Scientist | AI Enablement & Applied GenAI

Madison, NJ • 973-570-6331 • lazar.ryan123@gmail.com • linkedin.com/in/ryan-lazar123 • ryanlazar.dev

SUMMARY

Data scientist (M.S., May 2026) with 3 years at Deloitte translating ambiguous business problems into automated, AI-enabled solutions for federal and commercial clients. Hands-on with GenAI, LLM agents, and prompt engineering alongside core ML, with a track record of proof-of-concept tools and automation that cut manual effort and put decision-ready insight in front of non-technical stakeholders.

EXPERIENCE

Solution Analyst, Deloitte Consulting

Jul 2023 – Present

Philadelphia, PA

- Built and maintained Snowflake schemas across 100+ tables in dev and production, with automated ingestion pipelines serving daily reporting for state and commercial healthcare programs.
- Diagnosed and resolved recurring ETL failures from malformed data types, inconsistent delimiters, and null-handling errors; added validation logic that prevented downstream dashboard and compliance-reporting gaps.
- Automated extraction of hundreds of pages of Medicaid audit documentation into structured datasets with a custom Python script, replacing a manual process and delivering audit-ready data for a State of Georgia Medicaid review.
- Built a Python automation tool integrating REST APIs with SharePoint-connected Power BI dashboards, cutting manual data collection from 4 hours to 15 minutes daily (94%) and establishing a reusable team pattern.
- Maintained financial risk-monitoring data feeding 20+ Power BI dashboards for a \$42B federal broadband program, used by senior leadership to track spending and flag budget overruns.
- Translated complex state reporting mandates and payer quality standards into configured PCMS/PIMS scorecard logic, business rules, and system mappings, turning ambiguous requirements into a working solution.

Technical Project Manager, OneSeven Tech

Jul 2022 – Dec 2022

Miami, FL

- Led 12 people across 6 cross-functional teams (Design, Dev, QA), running sprint planning and velocity tracking that improved on-time delivery.
- Primary client liaison for 4–5 concurrent projects (\$45K–\$200K each), sustaining high satisfaction and renewal rates.

Executive Assistant to CFP, Northeast Financial Network

Nov 2021 – Jul 2022

Ledgewood, NJ

- Analyzed client portfolio performance and market data to support a CFP's planning and advisory work across 401(k), IRA, and insurance recommendations.

PROJECTS

AskMyData — Agentic Analytics on Oracle AI Database 26ai

Claude API · Oracle 26ai · FastAPI · Terraform · OCI

- Built and publicly deployed a conversational analytics agent with all AI infrastructure inside the database: in-database vector search (RAG), natural-language SQL over live data, and persistent agent memory — with read-only SQL guardrails and a full in-database audit trail.
- Provisioned OCI infrastructure with Terraform and shipped a containerized FastAPI service streaming the agent's reasoning loop live over SSE; served over HTTPS at askmydata.ryanlazar.dev.

NYC Taxi Analytics Platform

PySpark · Airflow · PostgreSQL · Next.js · Docker

- Engineered an Airflow-orchestrated PySpark ETL pipeline processing 100M+ NYC taxi records into PostgreSQL with data-quality checks and sub-second query performance, surfaced in a Dockerized Next.js/React dashboard for fare, demand, and hotspot analysis.

FHIR Patient Risk Stratification

Python · HL7 FHIR R4 · scikit-learn · Streamlit

- Built an end-to-end clinical ML pipeline ingesting Synthea FHIR R4 bundles, engineering 18 evidence-based features (SNOMED/ICD-10 comorbidities, vitals, utilization) to model 30-day readmission risk (0.96 AUC-ROC on synthetic Synthea data).
- Designed a Streamlit dashboard with cohort analytics and per-patient drill-down, styled to USWDS federal health standards; built on FHIR R4, the interoperability standard used by Epic.

Autonomous AI Investment Monitoring System

Anthropic API · Model Context Protocol (MCP) · Python · Supabase · Railway

- Built and deployed an autonomous two-agent system on the Anthropic API and Model Context Protocol that monitors holdings, pulls live market data via tool use, and generates daily risk and performance briefings with no human intervention.
- Implemented a grading engine computing beta, Sharpe ratio, and sector-diversification scores, with agents translating raw metrics into plain-language insights for non-technical decision-making.

TECHNICAL SKILLS

Languages: Python (pandas, NumPy, PySpark, scikit-learn), SQL, R

AI & GenAI: LLM application development, prompt engineering, AI agents (Anthropic API, Model Context Protocol), retrieval-augmented generation, GenAI workflow automation, Responsible AI

Data Engineering: ETL/ELT pipeline development, Apache Airflow, Informatica PowerCenter, Apache Spark, data-quality validation

Cloud & Warehousing: Snowflake, PostgreSQL, REST API integration, Docker

BI & Visualization: Power BI, Tableau, Streamlit

Machine Learning: Classification, regression, clustering, deep learning & NLP (TensorFlow/Keras), large language models

Tools & Practices: Git, JIRA, Agile / Scrum, stakeholder & change management, project management, data compliance

EDUCATION

M.S. Data Science, Drew University, Madison, NJ | GPA 3.76

May 2026

Coursework: Machine Learning, NLP & Deep Learning, Big Data Analysis, SQL & Relational Databases, Applied Regression

B.A. Economics, Minor in Data Science, Drew University, Madison, NJ | GPA 3.63

May 2023

A.S. Business Administration, Hudson County Community College, Jersey City, NJ | GPA 3.90

May 2021

CERTIFICATIONS Anthropic AI Fluency: Framework & Foundations • Anthropic AI Capabilities & Limitations • Microsoft PL-300 Power BI • Deloitte ML Foundation • PSM I (Scrum.org)